

Algorithmic choice relocates power

Research questions plus data artifacts collected (no results yet).

NO RESULTS TODAY


Algorithmic choice does not remove power - it relocates it.
Meeting focus: prioritize RQs and define what to recollect next.



System framing: pluggable services, new chokepoints



Takeaway: modular services relocate concentration.



Discovery surfaces and provider hosting are key chokepoints.



Bluesky UI: feed marketplace (/feeds)



Takeaway: feeds are discoverable + selectable algorithmic timelines.

A. /feeds directory

Search feeds

B. Feed cards

Creator + likes

C. Trending widget

Defaults matter

Discovery risk



Bluesky UI: what a feed looks like



Takeaway: a feed page renders an algorithm into a timeline of posts.

A. Feed name

Demand signals


B. Timeline output

Exposure decisions

C. Post actions

Why we measure

Manipulation risk



Bluesky UI: starter packs (onboarding bundles)



Takeaway: starter packs bundle people (and sometimes feeds) for fast onboarding.

A. Pack title

Curated list


B. Join button

1-click onboarding


C. People cards

Curation questions

Measure exposure



Research questions (R1-R5)



Takeaway: security and privacy scope.


R1  Discovery chokepoints? governance.

R2  Provider buckets dominate? reliability.

R3  Same winners across feeds? manipulation.


R4  Labels vary by viewer? safety.

R5  Future experiment: gaming plus privacy attacks. extra guardrails.





Data artifacts collected (no results)



Discovery inputs

feed_generators_index.csv
starterpacks.csv
starterpack_feeds.csv
discovery_feed_inclusions.csv
popular_feeds.csv

Panel snapshots

feed_panel.csv
feed_items.csv.gz
posts.csv.gz
authors.csv.gz
post_labels.csv.gz

Trust and reproducibility

state.db
run_metadata.csv
run_summary.csv
manifest.csv
data_dictionary.csv
validation_report.csv

Join spine

feed_panel -> feed_items -> posts/authors
post_labels attaches by feed_uri, viewer_mode, post_uri, post_cid.
Result: joinable, reproducible, auditable artifacts.



Data receipt: discovery surfaces (starter packs + popular list)



Takeaway: discovery is measurable via starter-pack inclusion and popular list ranks.

A. Starter packs

starterpacks +
starterpack_feeds

B. Inclusion_count

slot_count (position-
weighted)

C. Popularity_rank

Why it matters

Concentration + provider
leverage

A

starterpacks.csv (metadata)

name	creator_did	starterpack_uri
Interesting People	did:plc:32kt...u6sfyc	.../app.bsky.graph.starterpack/31r21bcjrd62

Receipt (truncated)

B

starterpack_feeds.csv (inclusion)

starterpack_uri	feed_uri
.../app.bsky.graph.starterpack/31r21bcjrd623	.../app.bsky.feed.generator/art-new

Receipt (truncated)

C

popular_feeds.csv (first row)

rank	feed_uri	collected_at_utc
1	.../app.bsky.feed.generator/whats-hot	2026-02-06T00:17:52Z

Receipt (truncated)

Data receipt: ranked exposure + risk labels



Takeaway: feed_items + posts + labels already support H1-H6 (cross-sectional).

A. feed_items row

(feed_uri, viewer_mode, rank)

feed_items.csv.gz (labeled impression)

feed_uri	viewer_mode	rank	post_uri
.../app.bsky.feed.generator/hodnauth	music	12	.../app.bsky.feed.post/3me5ms3r>g
post_cid	author_id		
bafyreidlca...t63ednqvh47a	did:plc:155j...1rqx7r		

B. posts row

text + links + langs

posts.csv.gz (metadata only)

post_uri	post_cid	author_id	created_at
.../app.bsky.feed.post/3me5ms3r>g	bafyreidlca...t63ednqvh47a	did:plc:155j...1rqx7r	2026-02-06T00:15:34.418Z
embed	external_domain	langs_json	
images	(none)	["en"]	

C. post_labels row

H6: risk exposure

Feed auditing / mitigations

post_labels.csv.gz (labels excerpt; Feb 1 run)

- did:plc:ar7c...vrndac:porn
- did:plc:227n...iy2skq:sexual
- did:plc:22ye...hw6pav:nudity
- did:plc:23rl...t4l754:graphic-media
- did:plc:ar7c...vrndac:self-harm
- did:plc:ar7c...vrndac:intolerant
- did:plc:ar7c...vrndac:rude
- did:plc:pa27...kcebyv:!no-unauthenticated
- did:plc:3dyb...4ws2vb:comment_count:0
- did:plc:37vi...zdr2m:reaction_count:5
- did:plc:3aqf...w4xz6v:Adult Content

Data receipt: labels (schema + global values + prefs)



Takeaway: labels are structured safety signals; behavior depends on viewer prefs.

A. Label schema

(src, uri/cid, val, neg, cts)

B. Global values

protocol-defined strings

C. Viewer prefs

hide / warn / ignore

Docs grounding

A

ATProto spec: label object fields

The fields on the label object are:

- ``ver`` (integer, required): label schema version. Current version is always ``1``.
- ``src`` (string, DID format, required): the authority (account) which generated this label
- ``uri`` (string, URI format, required): the content that this label applies to. For a specific record, an ``at://`` URI. For an account, the ``did:``.

B

Bluesky docs: global label values (excerpt)

- `!warn` which puts a generic warning on content but can be clicked through. Not configurable by the user.
- `!no-unauthenticated` which makes the content inaccessible to logged-out users in applications which respect the label.
- `porn` which puts a warning on images and can only be clicked through if the user is 18+ and has enabled adult content.
- `sexual` which behaves like `porn` but is meant to handle less intense sexual content.
- `graphic-media` which behaves like `porn` but is for violence / gore.
- `adult` which puts a warning on images but isn't 18+ and defaults to ignore.

C

Bluesky docs: viewer preferences (hide/warn/ignore)

```
moderationPrefs: {  
  // is adult content allowed?  
  adultContentEnabled: true,  
  
  // the global label settings (used on self-labels)  
  labels: {  
    porn: 'hide',  
  },  
}
```

Data receipt: custom labels (definitions + targets)



Takeaway: labelers can define custom taxonomies; we preserve raw label_val strings.

A. Custom definitions

blur + severity + defaults

A Bluesky docs: custom label values (table excerpt)

Blurs	Severity	Description
content	alert	Hide the content and put a "danger" warning label on the content if viewed
content	inform	Hide the content and put a "neutral" information label on the content if viewed

B. Label targets

account vs profile vs content

B Bluesky docs: label targets (account / profile / content)

Label targets

Labels may be placed on accounts or records. The Bluesky app interprets the targets in 3 different ways:

- On an account: has account-wide effects
- On a profile: affects only the profile record (e.g. user avatar) and never hides any content, including the account in listings.

C. Observed values

dataset examples

C Feb 1 run: custom label_val strings (examples)

```
- child_comment_count:2
- comment_count:5
- reaction_count:5
- Adult Content
- child_comment_count:0
- comment_count:0
- reaction_count:2
```

Why it matters

Data receipt: labelers (declarations + subscriptions + self-labels)



Takeaway: labels come from services; viewers choose which labelers to trust.

A. Declaration

values + policies

A

Bluesky docs: labeler declaration (record excerpt)

```
"type": "app.bsky.labeler.service",
"policies": {
  "labelValues": ["porn", "spider"],
  "labelValueDefinitions": [
    {
      "identifier": "spider",
      "severity": "alert",
      "blurs": "media",
      "defaultSetting": "warn"
```

B. Subscriptions

viewer-selected

B

Bluesky docs: labeler subscriptions (viewer-selected)

Labeler subscriptions

To include labels from a given labeler, you set the `atproto-accept-labelers` header on the HTTP request to a comma-separated list of the DIDs of the labelers you want to include. The AppView will subscribe to the included labelers (if it hasn't already) and attach any relevant labels. It includes the `atproto-content-labelers` header in the response to indicate which labelers were successfully included.

C. Self-labels

built-in values

Why it matters

C

Bluesky docs: self-labels (supported values)

Only global labels can be used for self-labeling, but not all global labels can be used for self-labeling. The supported self-label values are:

- `!no-unauthenticated`
- `porn`
- `sexual`

Data receipt: moderation decisions (flags + UI contexts)



Takeaway: the same label can blur/warn/filter depending on context + prefs.

A. moderatePost()

post + prefs → decision

B. Output flags

filter / blur / alert /
inform

C. UI contexts

contentList vs view vs
media

Docs grounding

A

Bluesky docs: moderatePost() (input → decision)

```
} from '@atproto/api'
```

Each of these follows the same API signature:

```
const res = moderatePost(post, moderationOptions)
```

B

Bluesky docs: decision flags (filter / blur / alert / inform)

The response object provides an API for figuring out what your UI should do in different contexts.

```
res.ui(context) /* =>
```

```
ModerationUI {  
  filter: boolean // should the content be removed from the interface?  
  blur: boolean // should the content be put behind a cover?
```

C

Bluesky docs: UI contexts (contentList vs contentView vs media)

- `displayName` The user's display name in any context
- `contentList` Content being listed, e.g. posts in a feed, posts as replies, a user list list, a feed generator list, etc
- `contentView` Content being viewed direct, e.g. an opened post, the user list page, the feedgen page, etc
- `contentMedia` Media inside the content, e.g. a picture embedded in a post

Here's how a post in a feed would use these tools to make a decision:

Collection credibility (one slide)



Read-only XRPC summary

Relay + AppView reads.

GET-only in this run.

POST auth only if auth mode is used.

No posting actions.

7-stage pipeline

relay -> index -> packs -> popular -> enrich -> panel
-> snapshot

Reproducibility: state.db, manifest.csv,
data_dictionary.csv, validation_report.csv.



What this enables next + advisor asks



Ready measurements

RQ1: measure discovery concentration.

RQ2: map provider leverage.

RQ3: test same winner overlap.

RQ4: compare label variability.

What I need from you

Ask 1: priority RQs and venue fit.

Ask 2: recollect longitudinal, auth mode, or experiments?

